Optimizing Fleet Efficiency: Recovering $33.5M in Revenue via Predictive Coordinated Routing

A data-driven approach to solving network-level congestion through geospatial analysis and a self-funding pilot model.

Jack Eisen

Table of Contents

Executive Summary:

Key Results at a Glance

- Total Revenue Leakage: **$113.7M (14.84%)** attributed to reactive routing inefficiencies.
- Strategic Recovery: **$33.5M** projected annual revenue at full scale (conservative **5% efficiency floor**).
- Capital Strategy: The pilot is designed to be **entirely self-funding**. By using targeted surcharges to cover initial costs, we can launch the initiative **without tapping into existing corporate budgets.**
- Financial Impact: **$18.5M Net Profit** in Year 1; **235% recurring ROI** by Year 2.

The Problem With Reactive Routing

Current fleet operations rely on individual GPS optimizations that fail at network scale. These systems trigger "herd-behavior"—simultaneously funneling drivers into the same alternative corridors after congestion is detected. This creates self-induced bottlenecks—a real-world manifestation of *Braess's Paradox*—where individual attempts to optimize routes actually degrade total network velocity and trap fleet earnings below the **median**.

The Solution to Reactive Routing:

This initiative proposes replacing reactive rerouting with a designed system that balances the fleet across the city. Key features **include proposed** *Anticipatory Distribution* to avoid "GPS stampedes," *Velocity Optimization* aligned with traffic signal timing, and *Real-time Synchronization* with external incident platforms.

Execution Strategy

Utilizing a **$1.64 RPM** (Revenue Per Minute) benchmark, the rollout follows two phases:

- **Phase 1 (Pilot):** Targeted deployment in **"Bottom 10"** performance segments. Funded by a **$0.99 surcharge** in high-performing **"funding hubs"** to achieve **$0 net cost.**
- **Phase 2 (Full-Scale):** Fleet-wide expansion to capture the identified **$33.5M recovery target.**

Market Inefficiency & Revenue Leakage

The Inefficiency Gap

Analysis reveals that predictable congestion is the primary ceiling on fleet growth. Even with standard GPS guidance, vehicles consistently hit bottlenecks that lower the average **RPM**. While individual drivers believe they are avoiding traffic, the lack of coordination means the fleet is often competing against itself for the same open pavement. External research (*Bilali et al., 2018*) indicates that transitioning to coordinated anticipatory routing can reduce travel time by up to **20%**, representing a massive untapped opportunity for efficiency gains. Standard GPS creates a 'selfish routing' environment; the Predictive Routing Initiative proposes a solution for *System Optimum*, mathematically neutralizing *Braess's Paradox* by preventing vehicles from over-concentrating on a single 'optimal' path—turning congestion into a manageable variable rather than an external constraint.

Establishing the $1.64 Performance Benchmark

Using **2**024 NYC TLC Trip Record Data, the median **RPM** across all trips was calculated to be **$1.64**. This represents the expected efficiency of trips operating under normal traffic conditions. The heavy right-tail distribution confirms that while the median is $1.64 RPM, a significant volume of trips operates at a deficit, representing the **$113.7M** leakage opportunity.

Figure 1: *Network-wide RPM Distribution showing the **$1.64** median benchmark used to identify the revenue leakage threshold.*

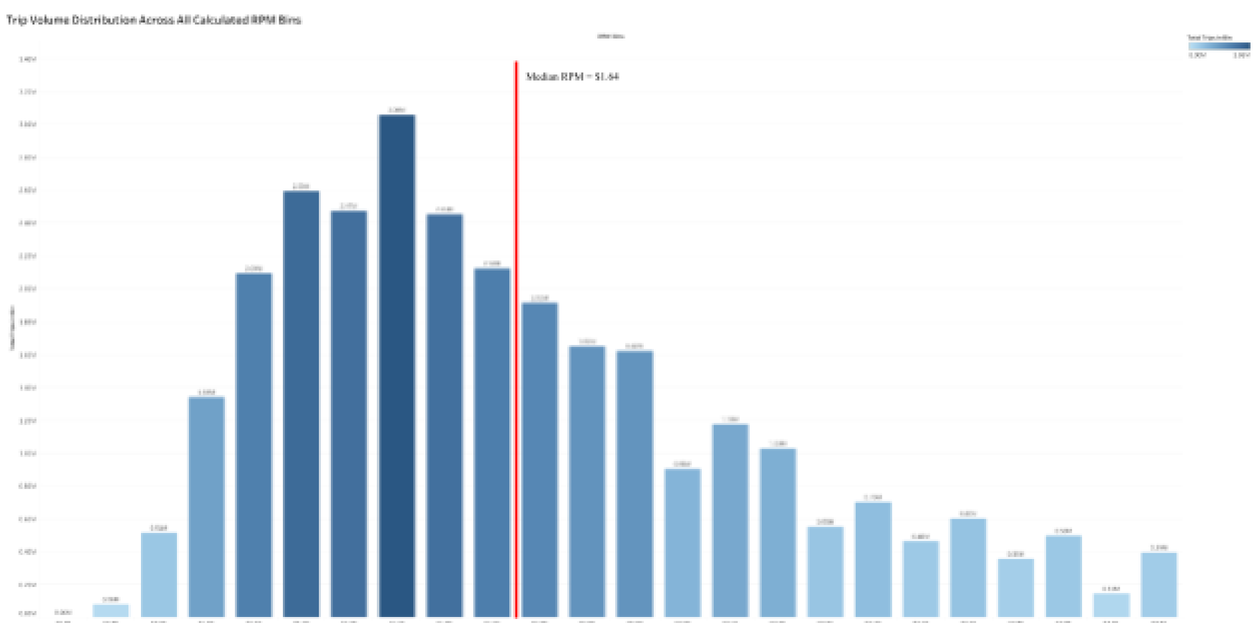

Key Insight: ***The heavy right-tail distribution*** *confirms that while the median is $1.64, a significant volume of trips operates at a deficit, representing the* ***$113.7M*** *leakage opportunity.*

Quantifying the Opportunity: The $113.7 Revenue Leakage Pool

Analysis indicates that trips falling below this threshold collectively account for a revenue leakage pool of **$113.7M**—revenue lost strictly due to extended trip durations. Because **RPM** is tethered to trip duration, **time** is our primary lever for monetization. By reducing avoidable delays through coordination, we increase revenue without needing to raise fares or increase trip volume.

Strategic Framework: Surgical Network Segmentation

Geospatial Analysis: Mapping Revenue Density via TLC Shapefiles

To move from theory to execution, a dual-filter of RPM and trip frequency was applied to the annualized dataset. These metrics were projected onto the **NYC TLC Taxi Zone Shapefile**, enabling a **spatial join** that segmented the network by Revenue Density. This process ensured that the identified 'Funding Hubs' and 'Recovery Zones' are grounded in official municipal boundaries rather than arbitrary clusters.

Figure 2: *Annualized trip metrics projected onto the* ***NYC TLC Taxi Zone Shapefile*** *to isolate geospatial revenue leakage and identify priority intervention zones.*

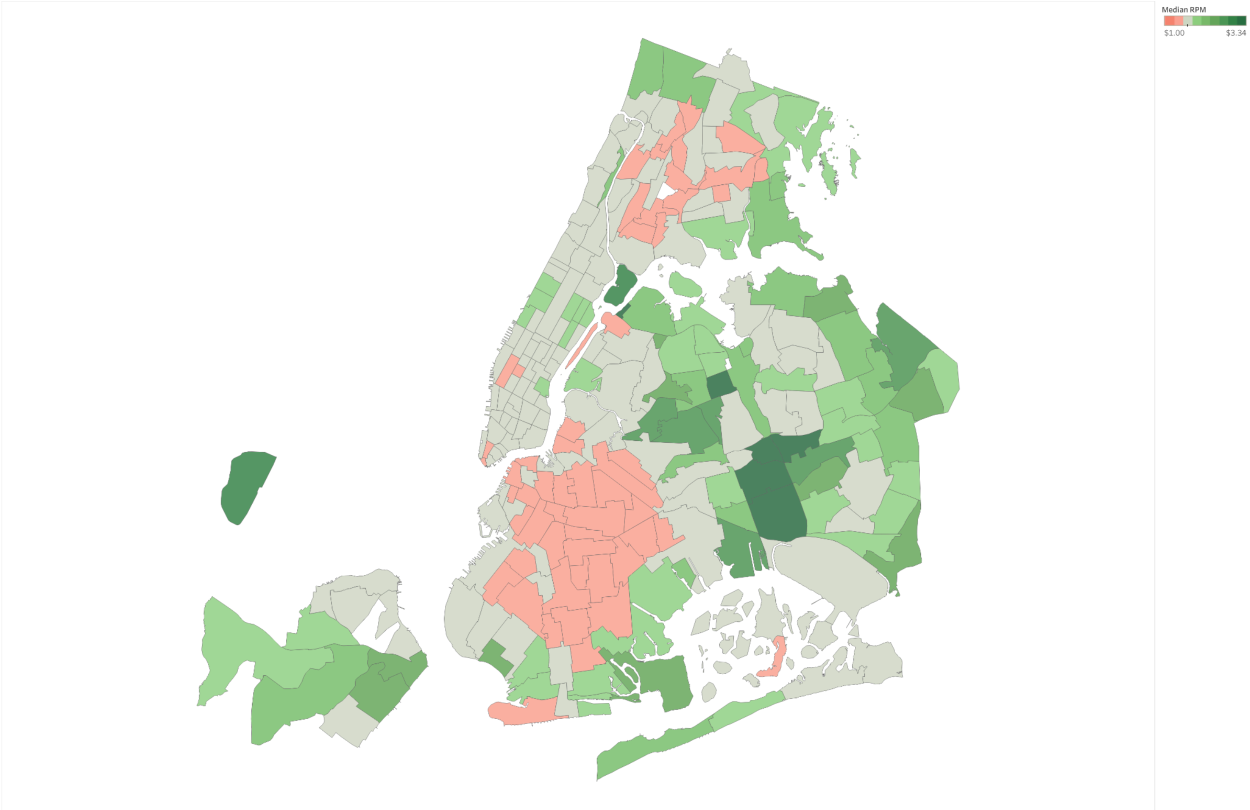

Operational Note: *Clusters in high-density corridors identify where* ***"GPS Stampedes"*** *are most frequent, providing the specific geospatial coordinates for the predictive routing layer.*

Segmentation Logic: Defining Funding Hubs and Recovery Zones

- **The Funding Source (Top 10 Segments):** High-demand corridors selected for consistent transaction volume. These zones are targeted for a **$0.99 strategic surcharge** to generate the **$1.13M pilot capital**.
- **The Performance Target (Bottom 10 Segments):** High-volume 'Cold Spots' identified via the TLC Taxi Zone Lookup Table and characterized by consistent duration variance. These serve as the testing ground to validate the **5% efficiency floor.**

Operational Baseline Audit: Establishing Pre-Intervention Metrics

A **'Value Beforehand'** baseline was established using **2024 audited revenue** for both the total fleet and the Bottom 10 segments. This ensures that the **5% efficiency gain** is applied to verified historical performance, removing the risk of over-projecting the **$33.5M recovery potential**.

Technical Architecture & Implementation Logic

Proposed Algorithmic Decision Layer: Solving Braess's Paradox

The proposed decision layer is designed to utilize *Constrained Optimization Algorithms* to solve *Braess's Paradox*. By calculating the **marginal cost** of adding one more vehicle to a specific corridor, the system prevents 'herd-behavior' from reaching the tipping point where rerouting degrades total network flow.

Proposed Infrastructure Requirements: OBD-II Telematics & IoT Integration

Proposed Infrastructure: **OBD-II telematics** will be utilized as the primary data source for real-time ingestion. This hardware is intended to enable the system to factor **engine idling** and **fuel consumption** into the **'Routing Brain's'** predictive model, increasing the accuracy of duration forecasting.

Data Integrity: NYC TLC 2024 Validation Methodology

The **$1.64 RPM** benchmark was derived using a relational join of the **NYC TLC 2024 Trip Record Data** and the Taxi Zone Lookup Table. Following professional data-clearing protocols, trips with durations under **1 minute** or fare amounts **≤ $0** were excluded. The NYC Taxi Zones Shapefile provided the geospatial basis for the spatial join, ensuring that the identified "Cold Spots" are grounded in official municipal boundaries.

Proposed Solution & Financial Engineering

Capital Strategy: The $0.99 Self-Liquidating Surcharge

To fund the proposed initial implementation, a temporary **$0.99 surcharge** will be applied to trips within the Top 10 performing segments. This provides the necessary transaction volume to generate **$1.13M in pilot capital**, offsetting the **$1.0M implementation cost**. This strategy offsets **100% of the initial hardware and training costs**.

Figure 3: *Analysis of **Top 10 High-Revenue Segments** identified as primary funding sources for the $1.13M pilot capital.*

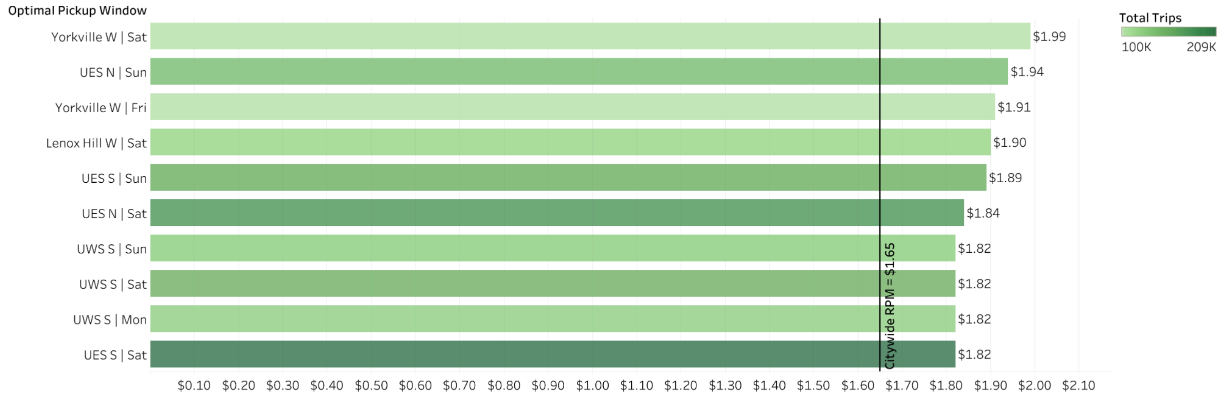

Pricing Elasticity: Validating the Low-Sensitivity Funding Model

The decision to implement a $0.99 Network Optimization Levy is supported by *Schaller (2021),* whose research on NYC for-hire vehicle demand indicates a low price-elasticity of approximately **-0.3** in high-density corridors. Because the "Funding Hubs" exhibit the lowest price sensitivity in the network, this **$1.13M capital** can be secured with negligible impact on trip volume.

Investment Lifecycle: From Pilot Surcharge to SaaS OpEx

The financial model transitions from a self-funded pilot to a traditional OpEx structure. While the initial $1.0M is offset by the strategic surcharge, Year 2 and onward will focus on fleet-wide scaling and the transition to a recurring SaaS maintenance model.

Table 1: *Multi-Year Investment & Funding Lifecycle*

Category	Pilot Phase (Bottom 10 Segments)	Year 1 (Fleet Scale Up)	Year 2+ (Recurring)
Funding Source	$0.99 Strategic Surcharge	General CapEx	OpEx
Hardware & IoT	OBD-II Telematics Retrofitting	Fleet wide Retrofitting	Replacement/Repairs
Software & AI	Centralized Decision Layer	Full API Integration	Cloud/SaaS Licensing
Operations	Driver Training (Pilot)	Fleetwide Onboarding	Ongoing Tech Support
Total Cost	**$1.0 Million (Offset by Surcharge)**	**$15.0 Million**	**$10.0 Million**

Valuation Logic: Benchmarking Against 2025 Enterprise Standards

The **$1.0M Pilot** and **$15.0M Year 1 expansion** are benchmarked against 2025 enterprise telematics standards. According to Tech.co (2025), professional-grade OBD-II hardware and AI-driven SaaS routing layers average **$300–$500 per vehicle** for initial deployment. Our $15M projection reflects a full-scale retrofit and centralized API integration for the entire active fleet, ensuring a **235% ROI** through the recovery of identified revenue leakage.

Projected Impact & ROI Analysis

The 5% Efficiency Floor: A Conservative Recovery Model

The Bottom 10 segments represent the primary testing grounds for the initiative. While these segments currently generate a combined **$37.44M baseline**, they suffer from the highest duration variance in the fleet. Applying a conservative **5% efficiency coefficient** to these specific zones projects an initial recovery value of **$1.96M**. While external research in coordinated routing (*e.g., Bilali et al., 2018*) suggests travel time reductions of up to **20%**, this model utilizes a **5% "Floor-First" approach**. This accounts for operational variables—including driver compliance and hardware latency—ensuring the **$33.5M revenue recovery** is a highly achievable target rather than an optimistic ceiling. This phase is designed to validate the system's ability to stabilize "Cold Spot" performance before scaling to the wider network.

Figure 4: *Performance Baseline for the* ***'Bottom 10'*** *pilot segments targeted for predictive routing optimization.*

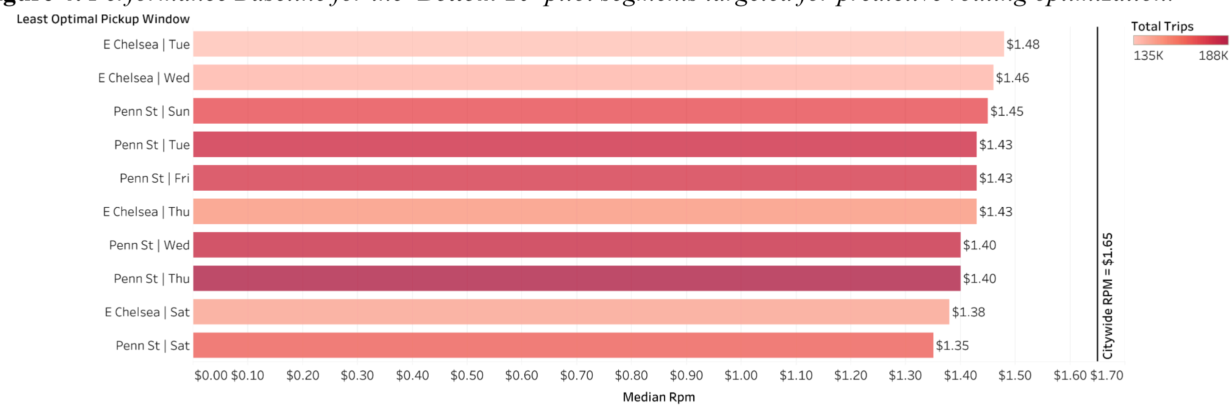

Target Logic: *These zones represent the network performance "floor" and are the primary targets for the* ***$1.96M pilot recovery phase.***

Full-Scale Network Impact: Achieving 235% Recurring ROI

Upon successful validation of the **5% gain** in the pilot group, the system will scale across the entire network. At this stage, the temporary **$0.99 surcharge is retired,** and the fleet transitions to a permanent predictive operational model. The focus shifts from capital generation to long-term value capture of the **$33.5M recovery target.**

Table 2: *Projected ROI and Network Impact (Fleet-Wide Scale)*

Financial Metric	Year 1 Projection	Year 2 Projection	Description
Net Capital Expenditure (CapEx)	$15 Million	$10 Million	Fleet wide hardwire/licensing (Yr 1); Maintenance (Yr 2+).
Incremental Revenue Recovered	$33.5 Million	$33.5 Million	Consistent revenue from the 5% efficiency gain.
Annual Net Profit	$18.50 Million	$23.50 Million	Revenue minus investment for the period
Annual ROI	**123.3%**	**235.0%**	Full capital recovery achieved within Year 1

Execution Roadmap: From Pilot to 235% Recurring ROI

Phase 1 & 2: Infrastructure & Pilot Funding (Months 1–9)

Deploy predictive routing hardware to **"Bottom 10"** vehicles and activate the **$0.99 strategic surcharge** in **"Top 10"** zones. This ensures the initial **$1.0M investment** is secured with a capital surplus and the pilot is fully operational with **zero net cost** to the general budget.

Phase 3: Audit & Surgical Scaling (Months 9–12)

Conduct a comprehensive audit of the **5% RPM gain** and driver compliance rates. Initiate fleet-wide API integration and synchronize the centralized "Routing Brain" with the broader network.

Phase 4: Full-Scale Expansion (Year 2+)

Retire the temporary surcharge once pilot capital targets are met. Execute the **$15M fleet-wide hardware retrofit** and transition to a SaaS-based maintenance model. According to *Global Market Insights (2025)*, the transition to AI-driven *Dynamic Route Optimization (DRO)* represents a critical shift in global logistics; our **$10.0M recurring OpEx** aligns with market valuations for high-tier predictive routing assets.

References

Bilali, A., & Allahviranloo, M. (2018). Analyzing the Impact of Anticipatory Vehicle Routing on the Network Performance. *Transportation Research Record*, 2672(20), 12–22. https://www.researchgate.net/publication/334094728_Analyzing_the_Impact_of_Anticipatory_Vehicle_Routing_on_the_Network_Performance

Braess, D. (1968). Über ein Paradoxon aus der Verkehrsplanung [On a Paradox of Traffic Planning]. *Unternehmensforschung*, 12, 258–268. https://link.springer.com/article/10.1007/BF01918335

Global Market Insights (2025). *Dynamic Route Optimization Software Market Size, Trends & Forecast 2025-2034.*https://www.gminsights.com/industry-analysis/dynamic-route-optimization-software-market

NYC Taxi & Limousine Commission (2024). *TLC Trip Record Data, Taxi Zone Lookup, and Shapefile.*https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page

Schaller, B. (2021). *Elasticities for Taxicab Fares and Service Availability.* Schaller Consulting. http://www.schallerconsult.com/taxi/elastic.pdf

Taale, H., & van Zuylen, H. J. (2001). The Impact of Anticipatory Route Guidance on Network Performance. *2001 IEEE Intelligent Transportation Systems Conference Proceedings.* https://ieeexplore.ieee.org/document/964353

Tech.co (2025). *Fleet Management Cost: In-Depth Explainer Guide.* https://tech.co/fleet-management/fleet-management-cost